

ShanghaiTech University

CS101 Algorithms and Data Structures

Fall 2025

Final Exam

Instructors: Yuyao Zhang and Xin Liu

Time: Jan 7th 8:00-10:00

INSTRUCTIONS

Please read and follow the following instructions:

- This exam has 8 questions, for a total of 100 points.
- You are not allowed to bring any papers, books, or electronic devices, including regular calculators.
- You are not allowed to discuss or share anything with others during the exam.
- You should write the answer to every problem in the dedicated box **clearly**.
- You should write **your name and your student ID** as indicated on the top of **each page** of the exam sheet.
- If you need more space, write “Continued on Page #” and continue your solution on the referenced scratch page at the end of the exam sheet.

Name	
Student ID	
Exam Classroom Number	
Seat Number	
(please copy this and sign)	<u>All the work on this exam is my own.</u>

THIS PAGE INTENTIONALLY LEFT BLANK.

DO NOT WRITE ANY ANSWER IN THIS PAGE!

1. (15 points) Single Choice

Each question has **exactly one** correct answer. Fill your answers **in the box below**.

Notice: Make sure to fully mark the answer, or we may take it unspecified.

(a)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(b)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(c)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(d)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(e)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D

- [3] (a) In a disjoint-set of size n that only uses path compression, what is the maximum change in height after a find operation with path compression?

- A. n .
- B. $n - 1$.
- C. **$n - 2$.**
- D. $n - 3$.

- [3] (b) Consider a set of precedence constraints among tasks labeled as $\{A, B, C, D, E, F, G\}$:

- A must be done before D and E ;
- B before D and F ;
- C before E and F ;
- D and E before G ;

Which order of completing tasks is **NOT** possible?

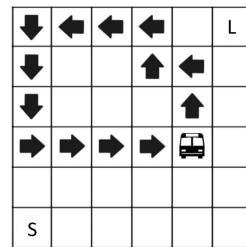
- A. A, B, C, D, E, F, G .
- B. B, C, A, E, D, F, G .
- C. C, B, A, D, E, F, G .
- D. A, B, C, E, G, D, F .**

- [3] (c) Suppose there is an undirected weighted connected graph with a very large number of vertices $|V|$, and the number of edges $|E|$ satisfies $|E| = \Theta(|V|^{1.5})$. All edge weights are non-negative. The task is to find the shortest paths from a single source vertex to all other vertices. Among the following combinations of algorithms and data structures, which one is optimal in terms of asymptotic time complexity?

- A. Using adjacency matrix storage with Dijkstra's algorithm that scans linearly to find the vertex with the minimum distance each time.
- B. Using adjacency list storage with Dijkstra's algorithm with a binary min-heap.**
- C. Using adjacency list storage with the Bellman-Ford algorithm.
- D. Using adjacency matrix storage with the Floyd-Warshall algorithm to compute all-pairs shortest paths.

- [3] (d) Logan needs to go to the school from his company every day. As shown in the figure, his company is located in the top-right corner (M, N) and the school is located in the bottom-left corner $(0, 0)$ of an $M \times N$ grid. He is able to move up, down, left, or right one grid per timestamp

on foot. However, there exists a subway route in the grid. He is allowed to take the subway along the route and move along the route at the rate of 3 squares per timestamp.



For A* search, consider the family of heuristics

$$h_k(x, y) = k \cdot (|x - 0| + |y - 0|),$$

where (x, y) is Logan's current location and the goal is $(0, 0)$.

What is the **largest** constant k such that h_k is guaranteed to be **consistent** for **any possible** subway route placement?

A. $k = 1$

B. $k = \frac{1}{2}$

C. $k = \frac{1}{3}$

D. $k = \frac{1}{4}$

- (e) We have a graph $G = (V, E)$ where $|V| = n$ and $|E| = m$. How many of the following problems have their **optimal** solution time complexity correctly labeled?
- (1) Determine whether there exists a path from a given vertex s to a given vertex t in an **unweighted directed** graph. $\Theta(n)$
 - (2) Given a **DAG**, compute a topological ordering of all vertices. $\Theta(n + m)$
 - (3) Find the shortest path distances from a source s to all vertices when edge weights are in $\{0, 1\}$. $\Theta((n + m) \log n)$
 - (4) Find the shortest path between **every pair** of vertices when $m = \Theta(n)$, edge weights $w \in \mathbb{N}^+$. $\Theta(n^3)$
 - (5) Determine whether there is a simple cycle in an undirected graph. $\Theta(n)$
- A. 1**
- B. 2
 - C. 3
 - D. 4

2. (20 points) Multiple Choices

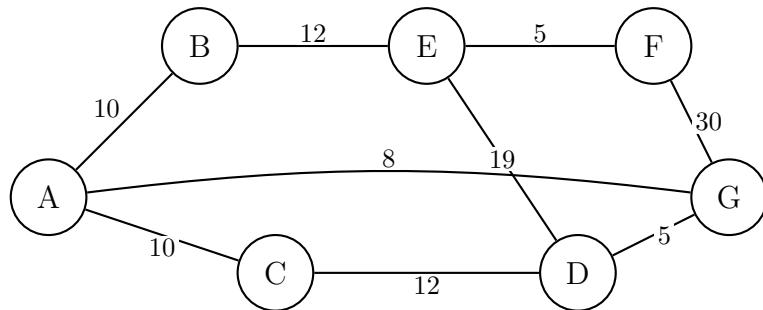
Each question has **one or more** correct answer(s). Select all the correct answer(s). For each question, you will get 0 points if you select one or more wrong answers, but you will get 2 points if you select a non-empty subset of the correct answers. Fill your answers **in the box below**.

Notice: Make sure to fully mark the answer, or we may take it unspecified.

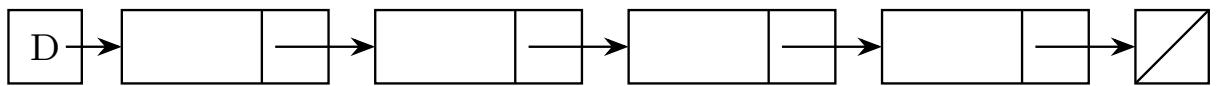
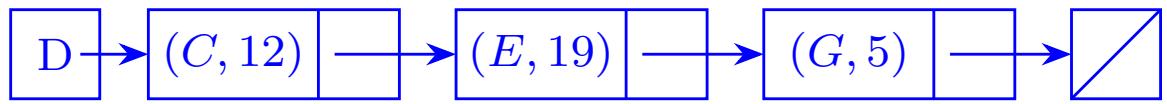
(a)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(b)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(c)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(d)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D
(e)	<input type="radio"/> A	<input type="radio"/> B	<input type="radio"/> C	<input type="radio"/> D

- [4] (a) In a weighted and connected graph $G = (V, E)$, which of the following statement(s) is/are **TRUE**?
- A. If two edges in E have the same weight, the graph must have multiple minimum spanning trees.
 - B. **If we add an edge to any spanning tree in G , there will be exactly one cycle in the new graph.**
 - C. The time complexity of Kruskal's algorithm is $O(|V| \log |E|)$.
 - D. Prim's algorithm can not work correctly when there are negative edges in the graph.
- [4] (b) Suppose we modify a graph algorithm to terminate early. Which of the following early-termination rules still guarantees the correctness of the algorithm's output?
- A. **Stop Kruskal's algorithm as soon as $V - 1$ edges have been added to the MST.**
 - B. **Assuming all edge weights are nonnegative, stop Bellman–Ford as soon as a full pass completes (all edges are relaxed) without decreasing any $\text{dist}[]$ value.**
 - C. Stop Floyd–Warshall if one outer-loop k produces no changes to the matrix.
 - D. Stop Dijkstra's algorithm as soon as every vertex has been inserted into the priority queue.
- [4] (c) Which of the following statement(s) is/are **TRUE**?
- A. **For an admissible heuristic function, A* tree search may have higher time complexity than A* graph search, but A* graph search may return a sub-optimal solution.**
 - B. In a connected undirected graph, regardless of the graph structure and implementation, Prim's and Kruskal's algorithms have the same time complexity.
 - C. If the graph is not connected, Floyd's algorithm cannot work correctly.
 - D. **In a directed graph, the Bellman–Ford algorithm can be used to detect negative cycles.**
- [4] (d) Which of the following statement(s) is/are **TRUE**?
- A. **In the Weighted Interval Scheduling problem with n intervals, both top-down (memoization) and bottom-up dynamic programming can achieve a time complexity of $O(n \log n)$.**

- B. In the *Coin Changing* problem with coin set $\{c_i\}_{i=1}^n$, let $OPT(v)$ be the minimum number of coins to make value v . Then $OPT(v) = \begin{cases} 0, & \text{if } v \leq 0, \\ \min_{1 \leq i \leq n} (OPT(v - c_i) + 1), & \text{if } v > 0. \end{cases}$
- C. From a DP perspective, in the graph $G = (V, E)$, the Floyd-Warshall algorithm defines $\Theta(|V|^2)$ sub-problems, and each sub-problem is solved in $\Theta(|V|)$ time, resulting in a time complexity of $\Theta(|V|^3)$.
- D. If there are n houses and m colors, the *House Coloring* problem can be solved in $\Theta(nm)$ time complexity.**
- [4]** (e) Suppose that Problem A is in P, Problem B is in NP, and Problem C is NP-Complete. Which of the following can you infer?
- A. **3-SAT polynomial time reduces to Problem C.**
- B. If Problem B can be solved in polynomial time, then P=NP.
- C. If Problem B cannot be solved in polynomial time, then neither can Problem C.**
- D. If Problem C polynomial time reduces to Problem B, then B is NP-Complete.**

3. (10 points) A World of Graph**I. Consider a weighted undirected graph $G = (V, E)$:**

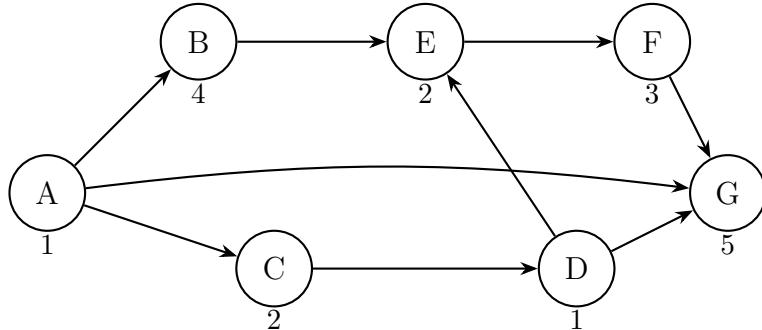
- 2 (a) For the weighted graph shown above, fill in the adjacency lists for vertex D . For a vertex u , its adjacency list is the set of all pairs (v, w) , where v is a neighbor of u and w is the weight of edge (u, v) . List the pairs in alphabetical order of v .

**Solution:**

- 1 (b) Fill in the blanks. For these questions, **Ties are broken by choosing the vertex whose label comes first alphabetically.**

- Is $\langle B, E, D, C, A \rangle$ a simple cycle? No
- Perform breadth-first search starting from vertex A . List the vertices in the order they are enqueued. $A \rightarrow B \rightarrow C \rightarrow G \rightarrow E \rightarrow D \rightarrow F$
- Perform recursive depth-first search starting from vertex A . List the vertices in the order they are visited. $A \rightarrow B \rightarrow E \rightarrow D \rightarrow C \rightarrow G \rightarrow F$
- Run Kruskal's algorithm on this graph. What are the edges in the MST? **Write their weights** in the order we add them. If the MST is not unique, write a valid possible solution. 5, 5, 8, 10, 10, 12

II. Now consider the directed graph below. The number below each vertex represents the vertex weight.



- 1 (c) Answer the questions (Write down all answers that satisfy the condition):

- [1] i. Which of the vertices have the same in-degree as out-degree? _____ **B, C, F**
- [1] ii. Which of the vertices have the minimum out-degree? _____ **G**
- (d) Now interpret each vertex as a task that can be executed in parallel, and let the processing time of each task be its vertex weight.
- [1] i. What's the **critical time of all tasks**? _____ **15**
- [1] ii. What's the corresponding **all critical path**? _____ **ABEFG**

4. (8 points) Trip

Alice is preparing for a trip. Alice wants to minimize her cost during the trip, which includes transportation and accommodation costs. There are $n + 1$ cities along the way. Alice starts at 0th city and the destination is the n^{th} city. The trip takes a few days, so if Alice stays overnight in the i^{th} ($1 \leq i \leq n$) city, she has to pay for the accommodation fee of a_i ($a_i > 0$). Alice must stay overnight at the destination. You are required to tell Alice the minimum cost for this trip, given the mode of transportation. Let $OPT(i)$ denote the minimum cost to arrive at city i and **stay overnight**.

- (a) If Alice decides to ride a bike, it will not incur any transportation costs. Alice can ride her bike through 1 or 2 cities within a day. That is, if Alice stays overnight at city i , Alice can stay overnight at city $i + 1$ or $i + 2$ the next day.

- 1 i. What is the answer to this question in terms of OPT ?

Solution: $OPT(n)$.

- 2 ii. Give your Bellman equation to solve the subproblems.

Solution:

$$OPT(i) = \begin{cases} a_i & \text{if } 1 \leq i \leq 2 \\ a_i + \min\{OPT(i-1), OPT(i-2)\} & \text{if } 2 < i \leq n \end{cases}$$

- (b) If Alice gives up riding a bicycle and decides to take a private jet, which can fly to any city but only once a day. Flying to city j from city i will incur $(i - j)^2$ cost.

- 3 i. Give your Bellman equation to solve the subproblems.

Solution: Define $OPT(0) = 0$.

$$OPT(i) = \min_{1 \leq j \leq i} \{OPT(i-j) + j^2\} + a_i$$

- 2 ii. What is the time complexity of your algorithm? (answer in $\Theta(\cdot)$ and in the most simplified form)

Solution: There are n sub-problems. For sub-problem i , we need to check all the sub-problems $j < i$. Therefore, the time complexity is $\Theta(n^2)$.

5. (11 points) Project Scheduling

As a student at HaishangTech, you are in RRR week. You have n projects due right now, but you haven't started any of them, so they are all going to be late. Each project requires d_i days to complete, and has a cost penalty of c_i per day. So if project i ends up being finished t days late, then it incurs a penalty of $c_i t$. Assume that once you start working on a project, you must work on it until you finish it, and that you cannot work on multiple projects at the same time.

For example, suppose you have three problem sets: CS100 takes 3 days and has a penalty of 12 points/day, CS101 takes 4 days and has a penalty of 20 points/day, and CS110 takes 2 days and has a penalty of 4 points/day. The best order is then CS101, CS100, CS110 which results in a penalty of $20 \times 4 + 12 \times (4 + 3) + 4 \times (3 + 4 + 2) = 200$ points.

- [3] (a) Describe your greedy algorithm that outputs an ordering of the projects that minimizes the total penalty for all the projects.

Solution: Schedule projects in nondecreasing order of the ratio $\frac{d_i}{c_i}$ (treat $c_i = 0$ as $\frac{d_i}{c_i} = +\infty$ and place them last). Output that ordering.

- [8] (b) Analyze the running time and prove the correctness of your algorithm by Exchange Arguments.

Solution: Compute ratios in $O(n)$ time and sort the n projects by $\frac{d_i}{c_i}$ in $O(n \log n)$ time. Total running time is $O(n \log n)$.

Simply: If unsorted, we can improve by swapping.

$$\begin{aligned} \frac{d_i}{c_i} > \frac{d_j}{c_j} \Rightarrow c_j d_i + (c_i d_i + c_j d_j) &> c_i d_j + (c_i d_i + c_j d_j) \\ \Rightarrow c_j(d_i + d_j) + c_i d_i &> c_i(d_i + d_j) + c_j d_j \end{aligned}$$

Here is the full proof:

1. Definition of a solution. A solution is an ordering (permutation) π of the n projects. If C_i is the completion time (days from now) of project i under π , then since all projects are due now, project i is finished $t = C_i$ days late and its penalty is $c_i C_i$. The total penalty is

$$P(\pi) = \sum_{i=1}^n c_i C_i.$$

2. Definition of optimality. An ordering π^* is optimal if it minimizes $P(\pi)$ over all orderings π .

3. How to change any optimal solution to a greedy solution iteratively. Take an optimal ordering π^* . If it is not sorted by nondecreasing $\frac{d_i}{c_i}$, then it contains an adjacent inverted pair: two consecutive projects i then j with

$$\frac{d_i}{c_i} > \frac{d_j}{c_j}.$$

Swap this adjacent pair. Repeat this operation while an adjacent inversion exists. When no adjacent inversion remains, the ordering is sorted by nondecreasing $\frac{d_i}{c_i}$, i.e., it matches the greedy rule.

4. Why does such a change not make the solution worse? Consider any adjacent pair i then j in a schedule, and let T be the total processing time of all projects before them (so T is fixed for this comparison).

If we do i then j , the pair contributes

$$P_{ij} = c_i(T + d_i) + c_j(T + d_i + d_j).$$

If we swap to j then i , the pair contributes

$$P_{ji} = c_j(T + d_j) + c_i(T + d_j + d_i).$$

Compute the difference:

$$P_{ij} - P_{ji} = c_j d_i - c_i d_j.$$

So $P_{ji} \leq P_{ij}$ (swapping is not worse) exactly when $c_j d_i \geq c_i d_j$, equivalently

$$\frac{d_i}{c_i} \geq \frac{d_j}{c_j}.$$

Therefore, whenever there is an adjacent inversion $\frac{d_i}{c_i} > \frac{d_j}{c_j}$, swapping that pair does not increase the total penalty (it decreases it). Hence, repeatedly swapping inversions never makes an optimal schedule worse, and the final inversion-free (ratio-sorted) schedule is also optimal. This is exactly the greedy schedule, so the greedy algorithm is correct.

6. (12 points) Dijkstra's Algorithm on Negative-weight Graphs

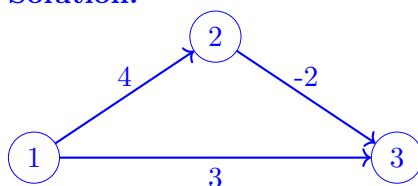
There's a directed graph $G = (V, E)$ without negative cycles. There could be negative-weight edges in E .

- [2] (a) We know that Dijkstra's algorithm doesn't work correctly on negative-weight graphs. Please give an example of G to show that Dijkstra's algorithm fails on G starting at vertex 1.

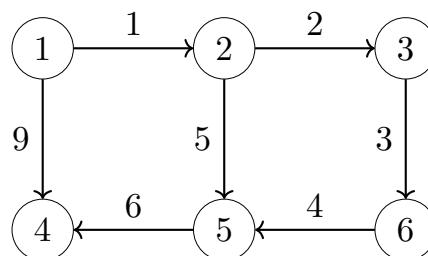
Requirements:

- $V = \{1, 2, 3\}$.
- There are no negative cycles in G .

Solution:



- (b) Given a graph G :



- [2] i. Please find the shortest path from vertex 1 and fill the distance in the table below.

Vertex	1	2	3	4	5	6
Distance	0	_____	_____	_____	_____	_____

- [2] ii. If we add a negative edge $(2, 6)$ with weight -8 . Please find the shortest path from vertex 1 and fill the distance in the table below.

Vertex	1	2	3	4	5	6
Distance	0	_____	_____	_____	_____	_____

Solution:

Vertex	1	2	3	4	5	6
Distance (i.)	0	1	3	9	6	6
Distance (ii.)	0	1	3	3	-3	-7

- [6] (c) If there is exactly **one** edge with negative weight in G , denoted as (p, q) . Please design an algorithm that is more efficient than Bellman-Ford (i.e., $O(|V||E|)$ time) to calculate the length of the shortest path from s to other vertices. You may assume $s, p, q \in V$ are distinct, and you can call Dijkstra's algorithm directly.

Hint: you can run multiple instances of Dijkstra's algorithm from different "sources".

Write your answer on the next page....

(c) Answer:

Solution: Algorithm:

1. Construct sub-graph $G' = (V, E')$, where $E' = E \setminus \{(p, q)\}$.
2. Run the Dijkstra's algorithm on G' from s, q , respectively, storing the distances in $d_s[\cdot], d_q[\cdot]$.
3. The shortest path from s to u must be one of the following situations:
 - The path not contains edge (p, q) , the distance of the shortest path in G is $d_s[u]$.
 - The path contains edge (p, q) (i.e. $[s, \dots, p, q, \dots, u]$), the distance of the shortest path in G is $d_s[p] + w(p, q) + d_q[u]$.

Therefore, in the original graph G , the length of shortest path from s to u is $\min\{d_s[u], d_s[p] + w(p, q) + d_q[u]\}$.

Time complexity analysis:

- Construct sub-graph: $O(1)$ (or $O(|V| + |E|)$) time.
- Run 2 times of Dijkstra's algorithm: $O(|E| \log |V|)$ (or $O(|V|^2)$) time.
- Calculate answers for every vertex: $O(|V|)$ time.

Therefore, the time complexity is $O(|E| \log |V|)$ (or $O(|V|^2)$).

7. (12 points) Marble Game

Logan and Joel are playing a strategy game on a directed acyclic graph (DAG) $G = (V, E)$ with n vertices and m edges. Each directed edge $(v \rightarrow u) \in E$ has an integer weight $w(v, u) \in \{1, 2, \dots, K\}$, where K is the maximum edge weight.

The game proceeds as follows:

- Logan's marble starts at vertex $s_L \in V$, and Joel's marble starts at $s_J \in V$.
- Logan moves first, and they alternate turns.
- On a player's turn, they move their marble along an outgoing edge $(x \rightarrow y) \in E$ such that $w(x, y) \geq \ell$, where ℓ is the weight of the edge used by the opponent in their last move (initialized to 0 before the first move).
- If a player cannot make a valid move, they lose.

Both players play optimally. Your task is to determine the winner of the game (Logan or Joel) for **every possible starting pair** $(s_L, s_J) \in V \times V$.

2 (a) How will you define the sub-problems?

Solution: Define $dp(a, b, k)$ as a boolean value representing whether the player whose turn it is can win from vertex a for the current player and vertex b for the opponent, with the last edge weight having code k .

6 (b) Describe and justify the Bellman Equation for computing the solution to sub-problems.

Solution: The recurrence for the DP is as follows:

- If there is no valid move from vertex a with a weight $\geq k$, then the current player loses.
- Otherwise, for each outgoing edge $(a \rightarrow a')$ with weight $w(a, a') \geq k$, the current player has two options:
 - Move to vertex a' with weight $w(a, a')$, and the opponent now has a turn to play.
 - If the opponent is forced into a losing position (i.e., $dp(b, a', w(a, a')) = \text{false}$), then the current player wins.

The recurrence relation is:

$$dp(a, b, k) = \begin{cases} \text{false} & \text{if there is no edge } (a \rightarrow a') \text{ with } w(a, a') \geq k, \\ \text{true} & \text{if there exists an edge } (a \rightarrow a') \text{ with } w(a, a') \geq k \\ & \text{and } dp(b, a', w(a, a')) = \text{false}, \\ \text{false} & \text{otherwise.} \end{cases}$$

Or

$$dp(a, b, k) = \bigvee_{\substack{(a \rightarrow a') \in E \\ w(a, a') \geq k}} \neg dp(b, a', w(a, a')),$$

This ensures that we always check if there exists a move that forces the opponent to lose, following the optimal play strategy.

Solution:

- [2] (c) What is the solution in terms of your sub-problems?

Solution: The solution to the original problem, i.e., the winner of the game for all initial positions of the marbles, is given by $dp(s_L, s_J, 0)$, where s_L is the starting position of Logan's marble, and s_J is the starting position of Joel's marble. If $dp(s_L, s_J, 0) = \text{true}$, Logan wins; otherwise, Joel wins.

- [2] (d) What is the runtime complexity of your algorithm? (answer in $\Theta(\cdot)$ and in the most simplified form, and give proof).

Solution: The number of sub-problems is $O(n^2k)$, and each fixed $dp(*, b, k)$ sub-problem requires $O(m)$ time to compute (since we check all outgoing edges of a vertex). Therefore, the total time complexity is $\Theta(nmk)$.

8. (12 points) Independent Set Partition

There is a **decision problem**, called Independent-Set-Partition ($\text{ISP}_{n,m}$):

In an undirected graph $G = (V, E)$ and two positive integers n and m , **determine** whether the vertex set V can be partitioned into exactly n subsets (some subsets may be empty) such that:

- Each subset is an *independent set* of G .
- The size of each subset is at most m .

The yes-instances $\text{ISP}_{n,m}$ is:

$$\text{ISP}_{n,m} = \left\{ \langle G = (V, E), n, m \rangle \mid \begin{array}{l} \text{In a simple undirected } G = (V, E), \exists V_1, \dots, V_n \subseteq V \text{ s.t.} \\ \text{- For every pair of } (V_i, V_j) : V_i \cap V_j = \emptyset, \cup_{i=1}^n V_i = V \\ \text{- For any } u, v \in V_i : (u, v) \notin E \text{ and } |V_i| \leq m. \end{array} \right\}$$

Note: Recall **k-Coloring**: Given an undirected graph $G = (V, E)$, **determine** whether there exists a coloring of the vertices with at most k colors such that no two adjacent vertices share the same color. Here is its yes-instance:

$$\text{k-Coloring} = \left\{ \langle G = (V, E), k \rangle \mid \begin{array}{l} G = (V, E) \text{ is an undirected graph, and there exists a coloring} \\ c : V \rightarrow 1, 2, \dots, k \text{ such that for every edge } u, v \in E, c(u) \neq c(v). \end{array} \right\}$$

- [2] (a) Prove that $\text{ISP}_{n,m}$ is in NP (Show your certificate and certifier with a brief explanation.)

Solution:

Certificate: A disjoint t -partition $\bar{V} = [V_1, \dots, V_t]$.

Certifier: First check whether \bar{V} is a disjoint partition and check each set $|V_i| \leq m$, which can be computed in $O(|V|)$ time. Then check whether for every V_i those vertices in it are independent, which takes $O(t|V|^2)$ time.

Given an input $s = \langle G = (V, E), n, m \rangle$, if there exists a certificate \bar{V} passes the certifier, then the V can be partitioned into \bar{V} , meaning $s \in \text{ISP}_{n,m}$. Otherwise, V can't be divided in that way, indicating $s \notin \text{ISP}_{n,m}$.

- [4] (b) If $n = 2$, the problems may be a little different from $n > 2$. For convenience, you can assume **the graph is connected**. From what you learned, show that $\text{ISP}_{2,m} \in \mathsf{P}$, more exactly, takes $O(|E| + |V|)$ time.

Solution: Use BFS on graphs: From a vertex $v \in V$, note that the two subsets are symmetry. W.L.O.G assume $v \in V_1$. Put v in the queue. Do till the queue is empty or conflict happens: Get the front of the queue v , then for any neighbor u of v :

- If u is visited, and u, v is in the same set. Then G is a no-instance of $\text{ISP}_{2,m}$.
- If u is visited, and u, v are in different sets. Then ignore u .
- Otherwise, put u in the opposite set of the set of v . Push u into the queue.

All vertices will be visited once. Then check of size of the set V_1, V_2 less than m to find whether it's valid. It takes $O(|E| + |V|)$ time.

The correctness depends on using BFS to solve 2-coloring problems.

(c) Now you are required to prove $\text{ISP}_{n,m}$ is in NP-Complete when $n \geq 3$.

- i. Write your polynomial time reduction f from k-Coloring to $\text{ISP}_{n,m}$.

Solution: Set $t = k$, $k \geq 3$ infers that t -COLORING is NP-Complete. Given an instance $\langle G, k \rangle$ of t -COLORING, we construct an instance of $\text{ISP}_{k,|V|}$: $\langle G, k, |V| \rangle$, using the same graph. This construction clearly takes polynomial time.

- ii. Prove that x is a yes-instance of k-Coloring $\Rightarrow f(x)$ is a yes-instance of $\text{ISP}_{n,m}$.

Solution: Suppose G has a valid t -coloring $c : V \rightarrow \{1, \dots, t\}$. Construct an partition as: $V_i = \{v \in V | c(v) = i\}$. V_1, \dots, V_k is an partition.

Correctness:

For each subset V_i , $V_i \leq |V|$, and $\forall (u, v) \in V_i$, we have $c(u) = c(v)$, which means $(u, v) \notin E$. $V_i \leq |V|$ and V_i is an independent set of G . Therefore, $\langle G, |V|, k \rangle$ is a yes-instance of $\text{ISP}_{n,m}$.

- iii. Prove: $f(x)$ is a yes-instance of $\text{ISP}_{n,m} \Rightarrow x$ is a yes-instance of k-Coloring.

Solution: Suppose the $\text{ISP}_{n,m}$ instance has a feasible partition V_1, \dots, V_k . Define a coloring $c(v) = \sum_{i=1}^t i[v \in V_i]$ i.e. $c(v) = i$ if and only if $v \in V_i$.

Correctness:

For each edge $(u, v) \in E$, assume $c(u) = c(v) = x$. That is, $u, v \in V_x$, which means subset V_x is not an independent set. It's a contradiction. Therefore, $c(u) \neq c(v)$. Hence, c is a valid t -coloring of G .

THIS PAGE INTENTIONALLY LEFT BLANK.

Fill the circle if you need to continue your solution on this page: Problem: _____